

Workbench architecture and flow diagram

START

Regulator query
the initial ask

“Clean up AML...”

User types into a chat interface what they want to focus on.

1 Interview

An **AI interviewer** helps the user sharpen the query into a precisely scoped purpose statement -- in one sentence -- which must be ratified and is then frozen. Everything is recorded and auditable.

SCOPE

Purpose statement
precisely constructed

Distill the current federal AML/CFT program-establishment obligations (the program / “pillar” rules) for U.S. depository institutions, money services businesses, and broker-dealers...

The **purpose statement** scopes and dictates everything else that happens.

2 Source discovery

An **AI discovery agent** searches for all authoritative sources relevant to the ratified purpose statement. Users are asked for feedback on the scope and can add/remove sources. (All actions are logged and auditable.) The goal is to include everything that contributes to the sludge. Every identified conflict, redundancy, gap, and downstream change is linked back to sources in the corpus.

PREPARATION

Source corpus
the legacy rulebook

The **corpus**: the accumulation of relevant statutes, regulations, guidance, manuals, FAQs, and enforcement actions -- with original source links -- comprising the de facto legacy rulebook in which sludge resides.

3 Distillation

An **AI distillation agent** reads everything and organizes it into a single document -- the big picture carefully reconciled and linked to every detail.

ONE DOCUMENT, READIED

Derived Blueprint

the mess, organized and readied

The **derived blueprint** allows users to see the forest and the trees in one place: every redundancy, contradiction, and gap, traced to the relevant provisions, machine-checked.

OUTCOME

Consolidated Rulebook
the new rulebook

The sludge-free set of statutes, regs, guidance, etc., highly granular and aligned with the target blueprint.

5 Alignment

An **AI alignment agent** identifies and recommends modifications to the source materials to achieve alignment with the target blueprint, all logged in a detailed **crosswalk**.

ONE DOCUMENT, CLEANED

Target Blueprint

cleaned, clear, and coherent

The **target blueprint** reflects all of the cleanup and redesign decisions, in a single doc. It provides a clean, coherent, target state, the north star to anchor reform efforts.

SCREENSHOT

An end-to-end workspace

Users control model selection, the scope and corpus, and every decision through a single interface

Rulebook Workbench M1 · advisory prototype

session \$0.2117 (12 calls) · lifetime \$96.1903 (1034 calls)

GUIDES

How to read the Workbench

About the workbench & the method

Reviewer's Guide — Refactor

Reviewer's Guide — Redesign

PROGRAMS

aml-program-rules-ref...

⋮

government-sponsor...

⋮

liquidity-redesign-p2

⋮

liquidity-refactor-p1

⋮

uk-derivatives-reporting

⋮

new-program-slug

+

Overview

Purpose

Corpus

Derived Blueprint

Refactor

Target Blueprint

Align

Setup

The whole project at a glance — what each step produced, links to open every document, and the log of human decisions. [Full reading guide ↗](#)

aml-program-rules-refactor-demo

refactor mode

Distill the current federal AML/CFT program-establishment obligations (the program/pillars' rules) for U.S. depository institutions, money services businesses, and broker-dealers (including the FINRA SRO layer) under FinCEN, the federal banking agencies, and the SEC as of a pinned 2026 date into a derived blueprint, defect register, and refactored target blueprint with operation trace for the Foundry working group, using statutes, regulations, SRO rules, official guidance, and examination manuals as authority and enforcement actions as evidence only, treating other covered sectors (casinos, mutual funds, insurance, etc.), OFAC/sanctions, fraud, and state AML/licensing as boundary regimes and SAR/CTR/CDD operational duties as coordination layer.

218 logged human decisions · \$51.09 model spend

✓ All phases complete — download the program package below.

Download program package

everything, organized — documents, decision log, corpus, transcripts

P0

Purpose ratified · refactor mode

Distill the current federal AML/CFT program-establishment obligations (the program/pillars' rules) for U.S. depository institutions, money services businesses,

P1

Corpus 42 sources · frozen · 42 acquired

manifest sha256:c89621b3a1c0dc

INTERVIEW

ANSWER IN THE BOX BELOW

Say hello to start — e.g. "I want to clean up the liquidity regime." The interviewer takes it from there, one question at a time.

Your answer...

Send

Mike Hsu

Program Owner

Light mode

RECORD — OVERVIEW

DECISIONS (218) ↓

DL-218

ratification

Ratify Target Blueprint: 170 operations finalized, 0 parked to Redesign Backlog

Mike Hsu (Program Owner) — Done

DL-217

disposition

accept RELATE-OBLIGATION (OP-197) as codify for defects-statute#7

? (Policy Reviewer) — agree

DL-216

disposition

accept RELATE-OBLIGATION (OP-196) as clarify for defects-statute#7

? (Policy Reviewer) — agree

DL-215

disposition

accept RESOLVE-CROSS-REFERENCE (OP-195) as fill_gap for defects-regulation_fincen#11

? (Policy Reviewer) — agree

DL-214

disposition

reject RELATE-OBLIGATION (OP-193) as unresolved for defects-guidance_fincen#6

? (Policy Reviewer) — agree

Tabs across the top follow the workflow described here — Purpose, Corpus, Derived Blueprint, Refactor, Target Blueprint, Align — and each one opens the document that phase produced.

In Setup, users can select which foundation model to use for each agentic task, including open weight models, with an OpenRouter API key.

The right-hand rail is the append-only record: every human disposition, who made it, and why. Model spend is metered in the header.

AML demo run · 42 sources · 218 logged decisions

Cleanup and policy change stay separate

Substantive changes are permitted only in redesign mode, and only when anchored to an adopted objective.

Everything is traceable

Every statement traces to its source; every human judgment records who decided, what, and why. A colleague today or the public later can check the work.

Humans are always in charge

Blueprints, registers, crosswalks and drafts are advisory working papers addressed to an Authority. Promulgation and repeal remain the Authority's acts. AI only assists, it never acts.